

da24c006

October 13, 2024

1 DAL Assignment 8


Submitted By : Harshil Pandey (DA24C006)

```
[7]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import matplotlib.colors as colors

from sklearn.datasets import make_circles
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.base import BaseEstimator, ClassifierMixin, clone
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.svm import LinearSVC
from sklearn.discriminant_analysis import LinearDiscriminantAnalysis

import seaborn as sns
sns.set_theme()

[8]: X, y=make_circles(n_samples=500, noise=0.1, random_state=42, factor=0.2)
y=2*y-1
X_train, X_test, y_train, y_test=train_test_split(X, y, random_state=42)
plt.scatter(X[:, 0], X[:, 1], c=y+2, cmap=colors.ListedColormap(['cyan', 'red']));
```



```
[3]: def plot_decision_boundary(model, X, y, ax=None, alpha=1, plot_points=True):
    '''
        This function plots the decision boundary of a classifier.

        #### Parameters

        `model` : This is the classifier for which the decision boundary has to be
        ↪ plotted.

        `X` : The data matrix used to plot the decision boundary. Each row is a
        ↪ data point, each column is a feature.

        `y` : The label vector for the data matrix. This will be used to color the
        ↪ boundary.

        `ax` : matplotlib.pyplot.Axes() instance. May be provided to draw on a
        ↪ specific axes object.

        `alpha` : This controls the transparency of the decision boundary.
```

```

    `plot_points`: If True, then the individual data points will also be
    ↳plotted according to their true label.

    #### Returns
    If `ax` is None then a new axes object is returned which contains the plot.
    ↳Otherwise the passed reference is returned.
    '''

    if ax is None:
        ax=plt.axes()
        points_x=np.linspace(X[:, 0].min(), X[:, 0].max(), 100)
        points_y=np.linspace(X[:, 1].min(), X[:, 1].max(), 100)
        mesh=np.meshgrid(points_x, points_y)
        df=pd.DataFrame(np.column_stack((mesh[0].flatten(), mesh[1].flatten()))))

        predictions=model.predict(df)
        ax.scatter(df[0], df[1], s=5, marker='.', c=predictions+2,
                   cmap=colors.ListedColormap(['blue', 'yellow']), alpha=alpha)
    if plot_points:
        ax.scatter(X[:, 0], X[:, 1], c=y+2, cmap=colors.ListedColormap(['cyan',
        ↳'orange'])), edgecolors='k')

    return ax

```

2 Task 1

```

[216]: class AdaBoost(BaseEstimator, ClassifierMixin):
    def __init__(self, estimator, learning_rate=0.5, n_estimators=10,
    ↳random_state=0):
        self.estimator=estimator # the learner that will be used
        self.learning_rate=learning_rate # the learning rate that will be used
        self.n_estimators=n_estimators # the number of max estimators to add in
        ↳the ensemble
        self.actual_iterations=0 # actual iterations took to fit the model
        self.estimators=[] # the list of trained estimators
        self.weights=[] # alphas, weight of each estimator
        self.random_state=random_state
        self._is_fitted=False

    def calc_error(self, y_true, y_pred, distribution):
        '''
        This function returns the error of the current iteration based on the
        ↳true labels, predicted labels and
        the current distribution.

        #### Parameters

```

```

        `y_true` : The true labels of classification.

        `y_pred` : The predicted labels of classification in the current_
        ↪ iteration.

        `distribution` : The current distribution of data points.
        '''

    return ((y_true != y_pred)*distribution).sum()

def calc_weight(self, error):
    '''
        This function calculates the weight of the current iteration based on_
        ↪ the error returned by `calc_error`.

        If error is 0, then 10 (a large value) is returned. If error is at_
        ↪ least half, then 0.01 (a small value) is returned).

        Otherwise the value is calculated using the usual formula.
        '''

    if error==0:
        return 10
    elif error>=0.5:
        return 0.01
    else:
        return self.learning_rate * np.log2((1-error)/error)

def updated_distribution(self, y_true, y_pred, distribution, new_alpha):
    '''
        This function calculates the updated distribution for the next_
        ↪ iteration using the true labels,

        predicted labels, distribution and the weight (alpha) of the current_
        ↪ iteration.

        ##### Parameters

        `y_true` : The true labels
        `y_pred` : The predicted labels
        `distribution` : The current distribution of the data points
        `new_alpha` : The weight (alpha) of the current distribution
        '''

    un_normalised=np.exp((-new_alpha)*y_true*y_pred)*distribution
    return un_normalised/un_normalised.sum()

def fit(self, X, y):
    # initialization steps

```

```

n=X.shape[0]
self.estimatedors=[]
self.weights=[]
self.actual_itations=0
distribution=np.array([1/n]*n) # initially the distribution is uniform

for i in range(self.n_estimators):
    new_estimator=clone(self.estimator)
    new_estimator.fit(X, y, sample_weight=distribution) # fit according
↳to the distribution
    predictions=new_estimator.predict(X)
    error=self.calc_error(y, predictions, distribution) # calculate the
↳error
    new_alpha=self.calc_weight(error) # calculate the weight

    self.estimatedors.append(new_estimator)
    self.weights.append(new_alpha)
    self.actual_itations+=1

    if error==0:
        break

    distribution=self.updated_distribution(y, predictions,
↳distribution, new_alpha) # update the distribution

self._is_fitted=True
return self

def predict(self, X):
    ensemble_predictions=np.zeros(X.shape[0])
    for i in range(self.actual_itations):
        ensemble_predictions+=self.weights[i]*self.estimatedors[i].predict(X)

    positive=(ensemble_predictions>=0)
    ensemble_predictions[positive]=1
    ensemble_predictions[~positive]=-1
    return ensemble_predictions

def plot_decision_boundary_pair(self, X, y, alpha=1):
    """
    This function plots the decision boundary for each intermediate
↳estimator and the whole ensemble.
    """
    fig, axes=plt.subplots(ncols=2)
    for i in range(self.actual_itations):
        points=False
        if i==self.actual_itations-1:

```

```

        points=True
        plot_decision_boundary(self.estimateds[i], X, y, ax=axes[0],
                               alpha=alpha, plot_points=points)
    plot_decision_boundary(self, X, y, ax=axes[1], alpha=0.4)
    axes[0].set_xlabel('$X_1$')
    axes[0].set_ylabel('$X_2$')
    axes[1].set_xlabel('$X_1$')
    axes[1].set_ylabel('$X_2$')
    return fig, axes

def __sklearn_is_fitted__(self):
    return self._is_fitted

```

```
[8]: dt=DecisionTreeClassifier(max_depth=1)
dt.fit(X_train, y_train)
```

```
[8]: DecisionTreeClassifier(max_depth=1)
```

```
[78]: clsf=AdaBoost(dt, n_estimators=1000)
clsf.fit(X_train, y_train)
```

```
[78]: AdaBoost(estimator=DecisionTreeClassifier(max_depth=1), n_estimators=1000)
```

```
[79]: clsf.actual_iterations
```

```
[79]: 1000
```

```
[80]: predictions=clsf.predict(X_train)
(predictions==y_train).mean()
```


```
[80]: np.float64(1.0)
```

```
[81]: predictions=clsf.predict(X_test)
(predictions==y_test).mean()
```

```
[81]: np.float64(1.0)
```

The Decision Boundary We can see that the model gives perfect accuracy and the combined decision boundary is composed of straight lines. This is because the underlying model is a Decision Tree which has a decision boundary of straight lines.

```
[82]: fig, axes=clsf.plot_decision_boundary_pair(X_train, y_train, alpha=0.4)
fig.set_size_inches(8.5, 4)
plt.tight_layout()
```



3 Task 2

```
[278]: def model_report(model):
    '''
    This function prints the model accuracy on the Training and Test Set
    '''
    print('Train Performance:')
    predictions=model.predict(X_train)
    score=np.mean((predictions==y_train))*100
    print('\tAccuracy Score: %.3f'%score)

    print('Test Performance:')
    predictions=model.predict(X_test)
    score=np.mean((predictions==y_test))*100
    print('\tAccuracy Score: %.3f'%score)
```

The following sections tunes the Boosting model with different underlying weak learners. In each section the best parameters are found using Grid Search and the Decision Boundary is visualised.

3.0.1 Decision Stump

```
[275]: gscv=GridSearchCV(estimator=AdaBoost(estimator=DecisionTreeClassifier(max_depth=1)),
    param_grid={'learning_rate': [0.4, 0.5, 0.6],
    'n_estimators': [20, 50, 100, 200, 250]},
    cv=5, refit=False, n_jobs=-1)
gscv.fit(X_train, y_train)
```

```
[275]: GridSearchCV(cv=5,
    estimator=AdaBoost(estimator=DecisionTreeClassifier(max_depth=1)),
```

```

n_jobs=-1,
param_grid={'learning_rate': [0.4, 0.5, 0.6],
            'n_estimators': [20, 50, 100, 200, 250]},
refit=False)

```

```

[276]: print('Best Params:' + repr(gscv.best_params_))
print('Best Score: %.3f'%gscv.best_score_)

```

```

Best Params: {'learning_rate': 0.4, 'n_estimators': 20}
Best Score: 0.989

```

```

[277]: clsf=AdaBoost(estimator=DecisionTreeClassifier(max_depth=1), **gscv.
        ↪best_params_)
        clsf.fit(X_train, y_train)

```

```

[277]: AdaBoost(estimator=DecisionTreeClassifier(max_depth=1), learning_rate=0.4,
                n_estimators=20)

```

```

[279]: model_report(clsf)

```

```


Train Performance:
    Accuracy Score: 100.000
Test Performance:
    Accuracy Score: 99.200

```

```

[281]: fig, axes=clsf.plot_decision_boundary_pair(X_train, y_train, alpha=0.4)
fig.set_size_inches(8.5, 4)
plt.tight_layout()

```



3.0.2 Decision Tree with max_depth=3

```
[282]: gscv=GridSearchCV(estimator=AdaBoost(estimator=DecisionTreeClassifier(max_depth=3)),  
                        param_grid={'learning_rate': [0.4, 0.5, 0.6],  
                                    'n_estimators': [20, 50, 100, 200, 250]},  
                        cv=5, refit=False, n_jobs=-1)  
gscv.fit(X_train, y_train)
```

```
[282]: GridSearchCV(cv=5,  
                  estimator=AdaBoost(estimator=DecisionTreeClassifier(max_depth=3)),  
                  n_jobs=-1,  
                  param_grid={'learning_rate': [0.4, 0.5, 0.6],  
                              'n_estimators': [20, 50, 100, 200, 250]},  
                  refit=False)
```

```
[283]: print('Best Params:' + repr(gscv.best_params_))  
print('Best Score: %.3f'%gscv.best_score_)
```

Best Params: {'learning_rate': 0.4, 'n_estimators': 200}
Best Score: 0.995


```
[284]: clsf=AdaBoost(estimator=DecisionTreeClassifier(max_depth=3), **gscv.  
                  ↪best_params_)  
clsf.fit(X_train, y_train)
```

```
[284]: AdaBoost(estimator=DecisionTreeClassifier(max_depth=3), learning_rate=0.4,  
              n_estimators=200)
```

```
[285]: model_report(clsf)
```

Train Performance:
 Accuracy Score: 100.000
Test Performance:
 Accuracy Score: 100.000

```
[286]: fig, axes=clsf.plot_decision_boundary_pair(X_train, y_train, alpha=0.4)  
fig.set_size_inches(8.5, 4)  
plt.tight_layout()
```



3.0.3 Linear SVM

```
[287]: gscv=GridSearchCV(estimator=AdaBoost(estimator=LinearSVC()),
                        param_grid={'learning_rate': [0.4, 0.5, 0.6],
                                    'n_estimators': [20, 50, 100, 200, 250]},
                        cv=5, refit=False, n_jobs=-1)
gscv.fit(X_train, y_train)
```

```
[287]: GridSearchCV(cv=5, estimator=AdaBoost(estimator=LinearSVC()), n_jobs=-1,
                  param_grid={'learning_rate': [0.4, 0.5, 0.6],
                              'n_estimators': [20, 50, 100, 200, 250]},
                  refit=False)
```

```
[288]: print('Best Params:' + repr(gscv.best_params_))
print('Best Score: %.3f'%gscv.best_score_)
```

Best Params: {'learning_rate': 0.4, 'n_estimators': 250}

Best Score: 0.949

```
[293]: clsf=AdaBoost(estimator=LinearSVC(), **gscv.best_params_)
clsf.fit(X_train, y_train)
```

```
[293]: AdaBoost(estimator=LinearSVC(), learning_rate=0.4, n_estimators=250)
```

```
[294]: model_report(clsf)
```


Train Performance:

Accuracy Score: 100.000

Test Performance:

Accuracy Score: 99.200

```
[295]: fig, axes=clsf.plot_decision_boundary_pair(X_train, y_train, alpha=0.1)
fig.set_size_inches(8.5, 4)
plt.tight_layout()
```



3.0.4 Logistic Regression

```
[296]: gscv=GridSearchCV(estimator=AdaBoost(estimator=LogisticRegression()),
                        param_grid={'learning_rate': [0.4, 0.5, 0.6],
                                    'n_estimators': [20, 50, 100, 200, 250]},
                        cv=5, refit=False, n_jobs=-1)
gscv.fit(X_train, y_train)
```

```
[296]: GridSearchCV(cv=5, estimator=AdaBoost(estimator=LogisticRegression()),
                  n_jobs=-1,
                  param_grid={'learning_rate': [0.4, 0.5, 0.6],
                              'n_estimators': [20, 50, 100, 200, 250]},
                  refit=False)
```

```
[297]: print('Best Params:' + repr(gscv.best_params_))
print('Best Score: %.3f'%gscv.best_score_)
```

```
Best Params: {'learning_rate': 0.4, 'n_estimators': 250}
Best Score: 1.000
```

```
[298]: clsf=AdaBoost(estimator=LogisticRegression(), **gscv.best_params_)
clsf.fit(X_train, y_train)
```

```
[298]: AdaBoost(estimator=LogisticRegression(), learning_rate=0.4, n_estimators=250)
```

```
[299]: model_report(clsf)
```

Train Performance:

Accuracy Score: 100.000

Test Performance:

Accuracy Score: 100.000

```
[301]: fig, axes=clsf.plot_decision_boundary_pair(X_train, y_train, alpha=0.1)
fig.set_size_inches(8.5, 4)
plt.tight_layout()
```

